

Exploratory Data Analysis - Tennis

December 1, 2020

Load in required packages:

```
[1]: import pandas as pd
import numpy as np
from scipy import stats

import seaborn as sns
import matplotlib.pyplot as plt
from IPython.display import set_matplotlib_formats

%matplotlib inline
set_matplotlib_formats('pdf', 'svg')
```

1 Dataset Summary

This project focuses on data from ATP matches in 2019 (obtained from https://github.com/JeffSackmann/tennis_atp). The dataset contains match statistics, including from both the match winner and loser, including age and height as well as aces, double faults, 1st serve and 2nd serve wins, and more.

The variables with integers (number of points/shots) for match stats will be converted to percentages to account for different match lengths. “Winner” is coded with the prefix “w” and “Loser” is coded with the prefix “l”.

The overall aim is to explore which match factors are related to **a) the likelihood of winning a match** and **b) the win strength** (or the % of total points won).

Read data:

```
[2]: tennis_df = pd.read_csv('atp_matches_2019.csv')
print('Raw data contains information on', tennis_df.shape[0], 'tennis matches,
      ↳with', tennis_df.shape[1], 'variables\n')
print('Columns =', tennis_df.columns.tolist())
```

Raw data contains information on 2781 tennis matches with 49 variables

```
Columns = ['tourney_id', 'tourney_name', 'surface', 'draw_size',
'tourney_level', 'tourney_date', 'match_num', 'winner_id', 'winner_seed',
'winner_entry', 'winner_name', 'winner_hand', 'winner_ht', 'winner_ioc',
```

```
'winner_age', 'loser_id', 'loser_seed', 'loser_entry', 'loser_name',
'loser_hand', 'loser_ht', 'loser_ioc', 'loser_age', 'score', 'best_of', 'round',
'minutes', 'w_ace', 'w_df', 'w_svpt', 'w_1stIn', 'w_1stWon', 'w_2ndWon',
'w_SvGms', 'w_bpSaved', 'w_bpFaced', 'l_ace', 'l_df', 'l_svpt', 'l_1stIn',
'l_1stWon', 'l_2ndWon', 'l_SvGms', 'l_bpSaved', 'l_bpFaced', 'winner_rank',
'winner_rank_points', 'loser_rank', 'loser_rank_points']
```

2 Plan for Data Exploration

To explore the data, I'm going to first filter by variables that I am interested in, rename them, and convert the data to long format so that each row refers to a player, not a match. Next, I will deal with missing values (NaNs) and plot histograms and boxplots of my variables to investigate their distribution and any potential outliers. Where needed, variables will be transformed to optimise model performance. Finally, I'm going to explore correlations between my variables with scatter plots and I will summarise their main attributes separately by match winners and losers.

3 Data Cleaning and Feature Engineering

3.1 Format variables of interest

Make sure that the data does not contain any duplicate rows:

```
[3]: n_dup = sum(tennis_df.duplicated())
     if n_dup == 0:
         print('No duplicate rows in dataset')
```

No duplicate rows in dataset

Filter by the variables of interest:

```
[4]: # select variables of interest
my_vars = ['winner_name', 'winner_ht', 'winner_age',
           'loser_name', 'loser_ht', 'loser_age',
           'w_svpt', 'w_ace', 'w_df', 'w_1stIn', 'w_1stWon', 'w_2ndWon',
           'l_svpt', 'l_ace', 'l_df', 'l_1stIn', 'l_1stWon', 'l_2ndWon']
tennis_df = tennis_df[my_vars]

# rename some variables for consistency and for long formatting later:
tennis_df.rename(columns={"winner_name": "name_w", "winner_ht": "height_w",
                          ↪ "winner_age": "age_w",
                          "w_svpt": "svpt_w", "w_ace": "ace_w", "w_df": "df_w",
                          "w_1stIn": "1stIn_w", "w_1stWon": "1stWon_w",
                          ↪ "w_2ndWon": "2ndWon_w",
                          "loser_name": "name_l", "loser_ht": "height_l",
                          ↪ "loser_age": "age_l",
                          "l_svpt": "svpt_l", "l_ace": "ace_l", "l_df": "df_l",
                          "l_1stIn": "1stIn_l", "l_1stWon": "1stWon_l",
                          ↪ "l_2ndWon": "2ndWon_l"},)
```

```

        inplace=True)
tennis_df.head()

```

```

[4]:
      name_w  height_w  age_w  name_l  height_l \
0    Kei Nishikori    178.0  29.004791  Daniil Medvedev    NaN
1  Daniil Medvedev    NaN  22.885695  Jo-Wilfried Tsonga    188.0
2    Kei Nishikori    178.0  29.004791    Jeremy Chardy    188.0
3  Jo-Wilfried Tsonga    188.0  33.705681    Alex De Minaur    NaN
4  Daniil Medvedev    NaN  22.885695    Milos Raonic    196.0

      age_l  svpt_w  ace_w  df_w  1stIn_w  1stWon_w  2ndWon_w  svpt_l  ace_l \
0  22.885695    77.0    3.0    3.0    44.0    31.0    17.0    100.0    8.0
1  33.705681    52.0   10.0    1.0    33.0    28.0    14.0    77.0   17.0
2  31.882272    47.0    2.0    2.0    33.0    26.0    9.0    46.0   10.0
3  19.868583    68.0   12.0    2.0    43.0    34.0   15.0    81.0    1.0
4  28.010951   105.0   12.0    3.0    68.0    48.0   25.0    94.0   29.0

      df_l  1stIn_l  1stWon_l  2ndWon_l
0    6.0    54.0    34.0    20.0
1    2.0    52.0    36.0    7.0
2    3.0    27.0    15.0    6.0
3    2.0    60.0    38.0    9.0
4    5.0    56.0    46.0   19.0

```

Convert integers to percentages using the number of service points per player (w/l_svpt):

```

[5]: # winner
tennis_df['ace_per_w'] = (tennis_df['ace_w'] / tennis_df['svpt_w'].tolist()) * 100
tennis_df['df_per_w'] = (tennis_df['df_w'] / tennis_df['svpt_w'].tolist()) * 100
tennis_df['1stIn_per_w'] = (tennis_df['1stIn_w'] / tennis_df['svpt_w'].
    ↳tolist()) * 100
tennis_df['1stWon_per_w'] = (tennis_df['1stWon_w'] / tennis_df['1stIn_w'].
    ↳tolist()) * 100

# loser
tennis_df['ace_per_l'] = (tennis_df['ace_l'] / tennis_df['svpt_l'].tolist()) * 100
tennis_df['df_per_l'] = (tennis_df['df_l'] / tennis_df['svpt_l'].tolist()) * 100
tennis_df['1stIn_per_l'] = (tennis_df['1stIn_l'] / tennis_df['svpt_l'].
    ↳tolist()) * 100
tennis_df['1stWon_per_l'] = (tennis_df['1stWon_l'] / tennis_df['1stIn_l'].
    ↳tolist()) * 100

```

Add variable for the percentage of total points won by the winner (and loser):

```
[6]: total_points = tennis_df['svpt_w'] + tennis_df['svpt_l']
# points won includes 1st and 2nd serve wins on their serve, and wins on other
↳ serve
w_points = tennis_df['1stWon_w'] + tennis_df['2ndWon_w'] + (tennis_df['svpt_l'] -
↳ tennis_df['1stWon_l'] - tennis_df['2ndWon_l'])
l_points = tennis_df['1stWon_l'] + tennis_df['2ndWon_l'] + (tennis_df['svpt_w'] -
↳ tennis_df['1stWon_w'] - tennis_df['2ndWon_w'])

tennis_df['pt_per_w'] = (w_points / total_points) * 100
tennis_df['pt_per_l'] = (l_points / total_points) * 100
```

Filter the df again to include only those variables I want to explore:

```
[7]: tennis_df.drop(['svpt_w', 'ace_w', 'df_w', '1stIn_w', '1stWon_w', '2ndWon_w',
                    'svpt_l', 'ace_l', 'df_l', '1stIn_l', '1stWon_l', '2ndWon_l'],
                    axis=1, inplace=True)
```

3.2 Convert to long format

To ensure that the data is in the most useful format for further exploration, I'm going to convert it to code won/lost as a variable, so that each row represents a player's stats and not the match stats:

```
[8]: # add unique id per match:
tennis_df["match_id"] = range(tennis_df.shape[0])

# long format
tennis_df_long = pd.wide_to_long(tennis_df,
                                stubnames=["name", "height", "age",
                                ↳
                                ↳ "ace_per", "df_per", "1stIn_per", "1stWon_per", "pt_per"],
                                i="match_id", j="outcome",
                                sep="_", suffix='\w+')
tennis_df_long.reset_index(inplace=True)
tennis_df_long.head()
```

```
[8]:
```

	match_id	outcome	name	height	age	ace_per	\
0	0	w	Kei Nishikori	178.0	29.004791	3.896104	
1	1	w	Daniil Medvedev	NaN	22.885695	19.230769	
2	2	w	Kei Nishikori	178.0	29.004791	4.255319	
3	3	w	Jo-Wilfried Tsonga	188.0	33.705681	17.647059	
4	4	w	Daniil Medvedev	NaN	22.885695	11.428571	

	df_per	1stIn_per	1stWon_per	pt_per
0	3.896104	57.142857	70.454545	53.107345
1	1.923077	63.461538	84.848485	58.914729
2	4.255319	70.212766	78.787879	64.516129

```
3  2.941176  63.235294  79.069767  55.704698
4  2.857143  64.761905  70.588235  51.256281
```

Variable key:

* match_id = unique code for the match (so shared by two rows - one for winner, one for loser)
 * outcome = whether the player won (w) or lost (l) * ace_per = percentage of service points that were aces * df_per = percentage of service points that were double faults * 1stIn_per = percentage of service points where the first serve was in * 1stWon_per = percentage of 1st serves where the point was won * pt_per = percentage of total match points that the player won

3.3 Remove missing values

```
[9]: # How many NaNs do we have in the data:
tennis_df_long.isna().sum()
```

```
[9]: match_id      0
outcome          0
name             0
height          2527
age              2
ace_per          204
df_per           204
1stIn_per        204
1stWon_per        204
pt_per           204
dtype: int64
```

The above numbers reveal that “height” has a lot of missing data, and so will be a particularly poor variable to include, so I’m going to drop that from the analysis:

```
[10]: tennis_df_long.drop(['height'], axis=1, inplace=True)
```

Luckily, there is a lot of data, so I’m going to simply remove any rows that contain a NaN for any variable:

```
[11]: tennis_df_long.dropna(inplace=True)
```

```
[12]: print('Data frame contains',len(tennis_df_long["match_id"].unique()),'matches_
↳with complete data, including',tennis_df_long.shape[0], 'players\n')
tennis_df_long.head()
```

Data frame contains 2679 matches with complete data, including 5358 players

```
[12]: match_id outcome      name      age      ace_per      df_per \
0         0         w      Kei Nishikori  29.004791  3.896104  3.896104
1         1         w      Daniil Medvedev  22.885695  19.230769  1.923077
2         2         w      Kei Nishikori  29.004791  4.255319  4.255319
```

3	3	w	Jo-Wilfried Tsonga	33.705681	17.647059	2.941176
4	4	w	Daniil Medvedev	22.885695	11.428571	2.857143

	1stIn_per	1stWon_per	pt_per
0	57.142857	70.454545	53.107345
1	63.461538	84.848485	58.914729
2	70.212766	78.787879	64.516129
3	63.235294	79.069767	55.704698
4	64.761905	70.588235	51.256281

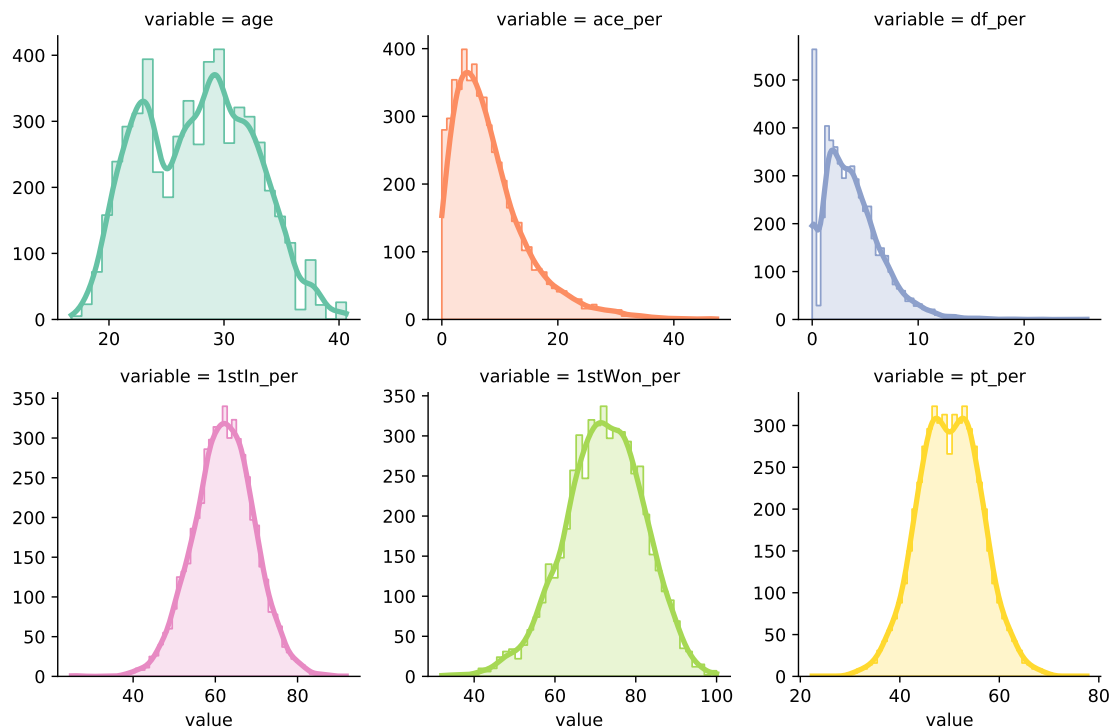
3.4 Distributions and data transformation

Now that the data is in a usable format, I'm going to inspect each variable to check for any outliers, or skewed distributions:

```
[13]: melted_df = tennis_df_long.select_dtypes("float").melt()

grid = sns.FacetGrid(melted_df, col='variable', col_wrap=3,
                      hue='variable', palette='Set2',
                      height=3, aspect=1,
                      sharex = False, sharey = False)
grid.map(sns.histplot, "value", alpha=.25,
         kde=True, line_kws={"lw":3}, element="step")
```

[13]: <seaborn.axisgrid.FacetGrid at 0x10a3a8fd0>



From the above plots, I can see that most of the variables are normally distributed. However, aces and double faults are substantially skewed, with most players making few of either in a match (a large number of players never hit a double fault). Therefore, I'm going to a) remove rows that have 0 aces or double faults and b) log-transform those variables:

```
[14]: tennis_df_long = tennis_df_long[(tennis_df_long["ace_per"] > 0) &
    ↳ (tennis_df_long["df_per"] > 0)].copy()
print('New data frame contains information for', tennis_df_long.
    ↳ shape[0], 'players')
```

New data frame contains information for 4567 players

```
[15]: # return the skew of my variables (> |.75| indicates skew that may need
    ↳ transforming)
skew_df = pd.DataFrame(stats.skew(tennis_df_long.select_dtypes("float")),
    index = tennis_df_long.select_dtypes("float").columns)
skew_df.rename(columns={0:"Skew"}, inplace=True)
skew_df.loc[skew_df["Skew"].abs() > .75, 'Transform'] = '*'
skew_df
```

```
[15]:
```

	Skew	Transform
age	0.104424	NaN
ace_per	1.543465	*
df_per	1.438449	*
1stIn_per	-0.088851	NaN
1stWon_per	-0.211778	NaN
pt_per	0.006292	NaN

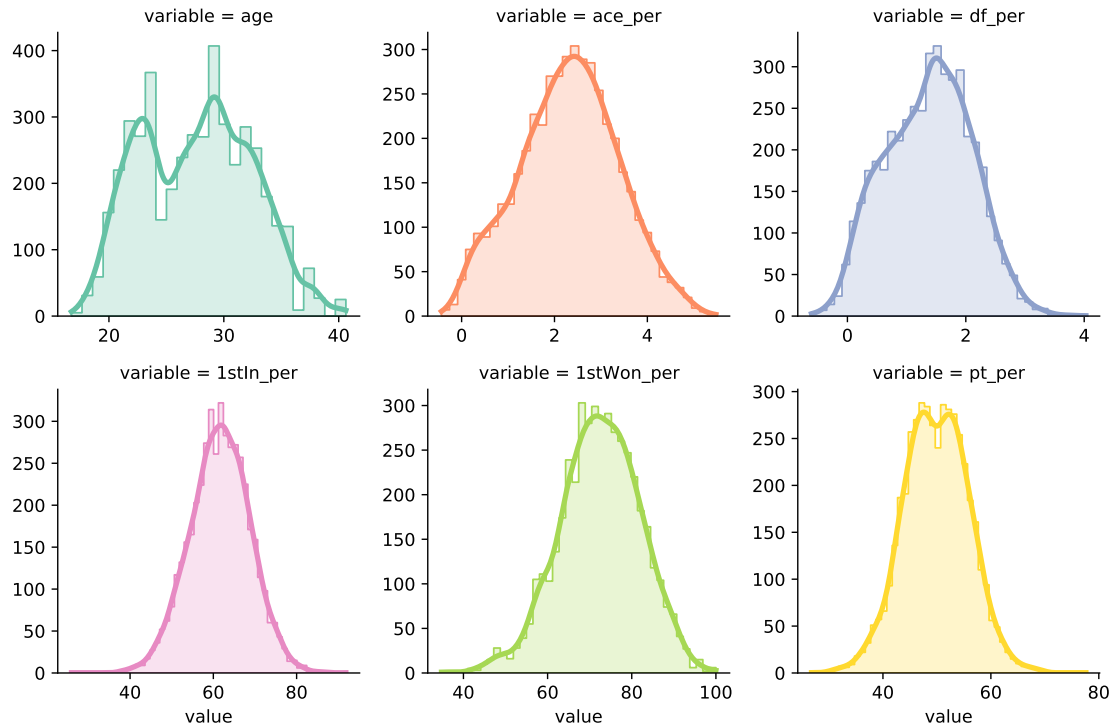
```
[16]: # log transform:
log_vars = skew_df.index[skew_df["Transform"]=="*"]
for v in log_vars:
    tennis_df_long[v] = stats.boxcox(tennis_df_long[v])[0]
```

Re-plot the log-transformed values:

```
[17]: melted_df = tennis_df_long.select_dtypes("float").melt()

grid = sns.FacetGrid(melted_df, col='variable', col_wrap=3,
    hue='variable', palette='Set2',
    height=3, aspect=1,
    sharex = False, sharey = False)
grid.map(sns.histplot, "value", alpha=.25,
    kde=True, line_kws={"lw":3}, element="step")
```

```
[17]: <seaborn.axisgrid.FacetGrid at 0x120922ba8>
```



Now all of the above variables appear to be normally distributed - great!


3.5 Outliers

Next, I am checking that there aren't any substantial outliers that might affect linear relationships. To do so, I'm first going to standardize my variables to put them all on an equivalent scale for plotting:

```
[18]: # standardize each of my variables to have equivalent scales:
tennis_df_z = tennis_df_long.copy()
tennis_df_z.loc[:,tennis_df_z.dtypes == "float"] = stats.zscore(tennis_df_z.
    ↪loc[:,tennis_df_z.dtypes == "float"], axis=0)
```

To be a bit more liberal with exclusion criteria (so marking fewer data points as outliers), I'm extending the whiskers for the below box plots to 2 * IQR (typically 1.5):

```
[19]: plt.figure(figsize=(6,5))
sns.boxplot(data=tennis_df_z.select_dtypes("float").melt(),
            y = 'variable', x = 'value', palette = "Set2",
            width=.6, whis=2)
plt.show()
```

Remove outliers from the data:

```
[20]: # data = a pandas series, marks outliers as NaN
def mark_outliers(data):
    Q1 = data.quantile(0.25)
    Q3 = data.quantile(0.75)
    IQR = Q3 - Q1

    my_filter = (data < Q1 - 2 * IQR) | (data > Q3 + 2 * IQR)
    return my_filter

[21]: my_vars = tennis_df_long.select_dtypes("float").columns.tolist()
for v in my_vars:
    filt = mark_outliers(tennis_df_long[v])
    tennis_df_long.loc[filt,v] = np.NaN

[22]: # drop missing values (outliers):
tennis_df_long.dropna(inplace=True)

print('Cleaned data contains statistics for',tennis_df_long.
      ↪shape[0], 'players\n')
```

```
tennis_df_long.head()
```

Cleaned data contains statistics for 4556 players

```
[22]:
```

	match_id	outcome	name	age	ace_per	df_per	\
0	0	w	Kei Nishikori	29.004791	1.536686	1.484146	
1	1	w	Daniil Medvedev	22.885695	3.879314	0.681774	
2	2	w	Kei Nishikori	29.004791	1.649625	1.589506	
3	3	w	Jo-Wilfried Tsonga	33.705681	3.735738	1.155998	
4	4	w	Daniil Medvedev	22.885695	3.042286	1.122826	

	1stIn_per	1stWon_per	pt_per
0	57.142857	70.454545	53.107345
1	63.461538	84.848485	58.914729
2	70.212766	78.787879	64.516129
3	63.235294	79.069767	55.704698
4	64.761905	70.588235	51.256281

4 Key Findings and Insights

In this section, I am using the cleaned data to provide some preliminary insight into how my player variables relate to the likelihood of winning a match as well as the total percentage of points won in a game.

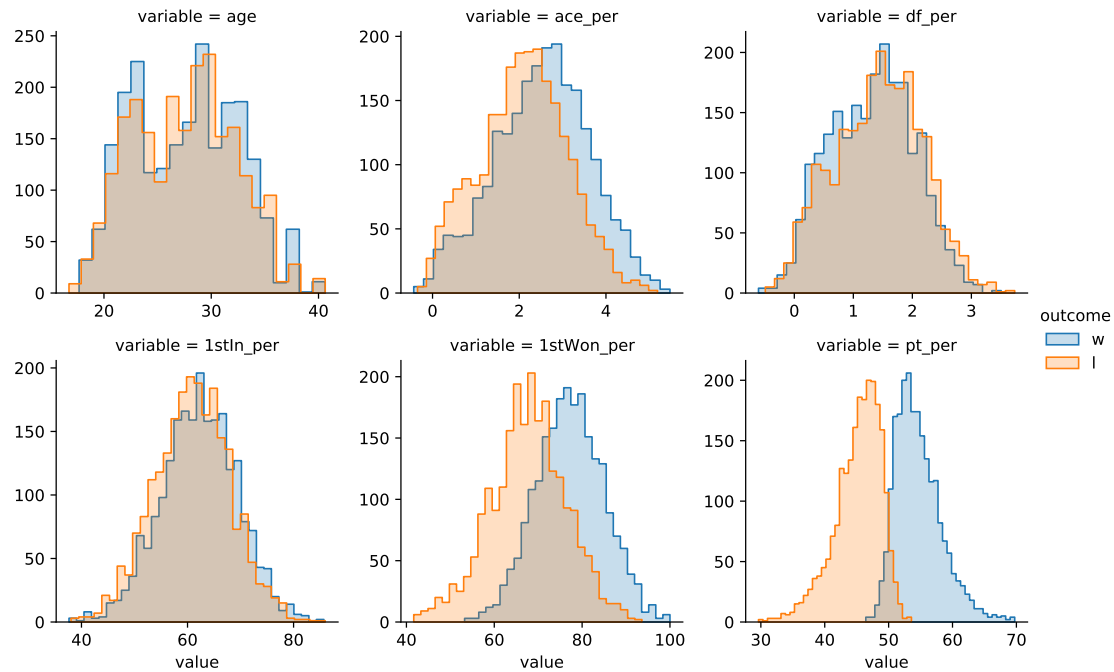
4.1 Data for winners vs. losers

Here are the same histograms as I plotted above, but now divided by whether the player won or lost their match:

```
[23]: melted_df = tennis_df_long.drop(['match_id', 'name'], axis=1).
      ↪ melt(id_vars="outcome")

grid = sns.FacetGrid(melted_df, col='variable', col_wrap=3,
                     hue='outcome',
                     height=3, aspect=1,
                     sharex = False, sharey = False)
grid.map(sns.histplot, "value", alpha=.25,
         element="step")
grid.add_legend()
```

```
[23]: <seaborn.axisgrid.FacetGrid at 0x120e889e8>
```



Unsurprisingly, players who win the match tend to win a higher percentage of total points (“pt_per”). Visually, what appears to be most strongly associated with the probability of winning is the percentage of aces made in a match, as well as the percentage of 1st service points won. Double faults and percentage of 1st serves in appear to make a small contribution, and there appears to be minimal or no impact of age.

Here is a summary of descriptive statistics, grouped by winners and losers:

```
[24]: tennis_df_long.drop('match_id', axis=1).groupby('outcome').
      ↪agg(['mean', 'std', 'min', 'max']).transpose()
```

```
[24]: outcome      1      w
age      mean  27.551390  27.672662
        std    4.710009   4.912380
        min   16.739220  17.681040
        max   40.624230  40.624230
ace_per  mean    2.098201   2.528593
        std    1.011213   1.086299
        min   -0.346641  -0.433622
        max    5.189438   5.476191
df_per   mean    1.424274   1.319377
        std    0.741199   0.708457
        min   -0.490891  -0.606692
        max    3.734195   3.504947
1stin_per mean   60.900438  62.281636
        std    7.130305   7.198408
```

	min	38.028169	37.647059
	max	85.964912	85.714286
1stWon_per	mean	67.741135	77.268828
	std	8.667596	7.701484
	min	41.666667	53.246753
	max	93.617021	100.000000
pt_per	mean	45.207064	54.683643
	std	3.666534	3.628172
	min	29.629630	46.428571
	max	53.571429	69.736842

4.2 Correlations


Due to the amount of data, I'm randomly selecting a sub-sample to plot the correlations between my variables:

```
[25]: print('Randomly sampling data of 150 players\n')
data_sample = tennis_df_long.sample(150, replace=False)

sns.pairplot(data_sample.select_dtypes("float"),
              kind="reg", height=1.4, corner=True,
              plot_kws={'scatter_kws':{'s':8},
                        'line_kws':{'color':'black',
                                    'lw':2}})

plt.show()
```

Randomly sampling data of 150 players



A couple of things to take away from the above plot:

- none of our feature variables (all apart from pt_per — percentage of game points won) show a particularly strong correlation with one another, so no real issues with multicollinearity.
- looking at the bottom row, we confirm the observations above comparing winners and losers — a higher percentage of game points won (pt_per) is positively associated with percentage of 1st serves won, as well as slightly correlated with aces and percentage of 1st serves in, and is slightly negatively associated with the double faults.

5 Potential Hypotheses

1. Match winners make a higher percentage of 1st serves than match losers.
2. Match winners are younger than match losers.

3. A higher number of match points won is associated with a higher percentage of 1st serves in.

6 Significance Test: Hypothesis 1

To test the hypothesis that match winners make a higher percentage of 1st serves than match losers, I'm running an independent samples t-test with scipy:

```
[26]: grp1 = tennis_df_long.loc[tennis_df_long["outcome"] == "w", "1stIn_per"]
      grp2 = tennis_df_long.loc[tennis_df_long["outcome"] == "l", "1stIn_per"]
      stats.ttest_ind(grp1, grp2)
```

```
[26]: Ttest_indResult(statistic=6.506143450550752, pvalue=8.541338317664799e-11)
```

A p value of $< .001$ shows that winners make a significantly higher % of 1st serves (at $\alpha = .05$) than losers.

7 Next Steps

The relationship between percentage of match points won and my continuous features — 1st serves, aces, double faults, etc. — could also be explored separately for winners and losers or based on age categories. For example, perhaps older players benefit (in terms of winning likelihood) from a higher number of aces than younger players? Additionally, you could use classification, such as logistic regression, to predict whether a player will win or lose based on their match statistics. In this case, you could train on half of the data and test on the other half of the data to test the generalisability of the model. Finally, I would be interested in looking at the differences between opponents (by match) rather than each player's individual performance — for example, perhaps age becomes important for winning likelihood only when there is a large difference in age between opponents.

8 Dataset Quality Summary

Overall, this dataset appears to be good quality. The only variable that had a significant amount of missing data was players' height, which was therefore excluded. Most of the variables are normally distributed, and relatively few outliers appeared in the data, and there do not appear to be redundancies between the variables (none show a very high correlation). Some additional stats would be useful, including forehand and backhand winners, rally durations, etc., in order to build a comprehensive model of winning likelihood.