

Survival_analysis_project

March 24, 2021

1 Survival Analysis Project: Customer Churn

1.1 Data Summary

The data analyzed here contains information about customer churn for a telecommunications company. The data can be downloaded [here](#).

The data includes a number of features about each user, such as their contract type, whether they stream movies and tv, their phone service, and so on.

1.2 Analysis Objectives

The goal of this analysis is to model which features best explain the probability of a customer churning in the future.

```
[1]: # load packages:
import pandas as pd
import numpy as np
from scipy.stats import zscore
import matplotlib.pyplot as plt
import seaborn as sns
from lifelines import KaplanMeierFitter, CoxPHFitter

# visualization
import matplotlib.pyplot as plt
import seaborn as sns
from IPython.display import set_matplotlib_formats

%matplotlib inline

# set to show vector images
set_matplotlib_formats('pdf', 'svg')
```

1.3 Data Overview

```
[2]: # load in the data:
data = pd.read_csv('telco-customer_churn.csv')
data.shape
```

```
[2]: (7043, 21)
```

```
[3]: # what are our features?
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customerID            7043 non-null   object
 1   gender                7043 non-null   object
 2   SeniorCitizen         7043 non-null   int64
 3   Partner               7043 non-null   object
 4   Dependents            7043 non-null   object
 5   tenure               7043 non-null   int64
 6   PhoneService          7043 non-null   object
 7   MultipleLines         7043 non-null   object
 8   InternetService       7043 non-null   object
 9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

First, I'm going to do some basic data cleaning and visualisation:

```
[4]: # check for any duplicates:
data.duplicated().sum()
```

```
[4]: 0
```

```
[5]: data.head()
```

```
[5]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female              0     Yes           No         1           No
1  5575-GNVDE   Male              0     No           No        34           Yes
2  3668-QPYBK   Male              0     No           No         2           Yes
3  7795-CFOCW   Male              0     No           No        45           No
```

4	9237-HQITU	Female	0	No	No	2	Yes
---	------------	--------	---	----	----	---	-----

	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	\
0	No phone service	DSL	No	...	No	
1	No	DSL	Yes	...	Yes	
2	No	DSL	Yes	...	No	
3	No phone service	DSL	Yes	...	Yes	
4	No	Fiber optic	No	...	No	

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

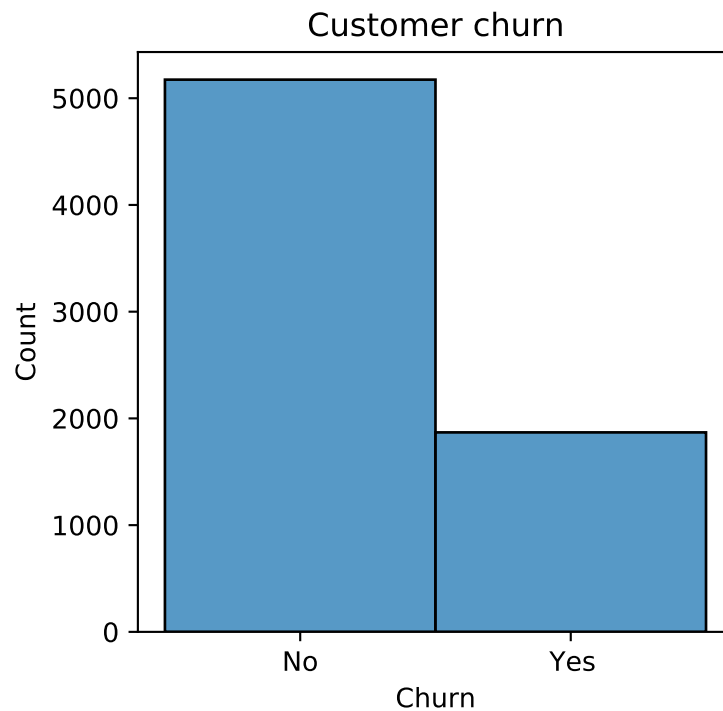
[5 rows x 21 columns]

```
[6]: # any missing values?
data.isna().sum()
```

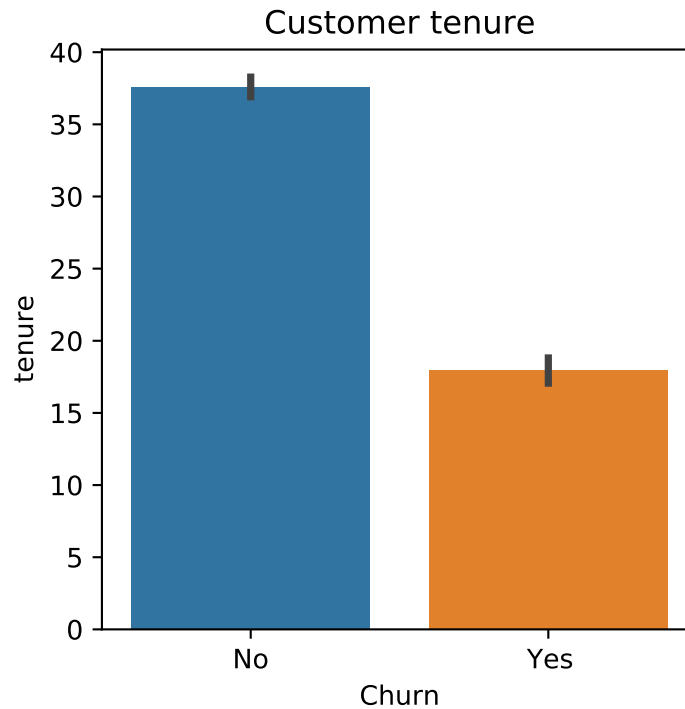
```
[6]: customerID      0
gender              0
SeniorCitizen      0
Partner            0
Dependents         0
tenure             0
PhoneService       0
MultipleLines      0
InternetService    0
OnlineSecurity     0
OnlineBackup       0
DeviceProtection   0
TechSupport        0
StreamingTV        0
StreamingMovies    0
Contract           0
PaperlessBilling   0
PaymentMethod      0
MonthlyCharges     0
```

```
TotalCharges      0
Churn              0
dtype: int64
```

```
[7]: # How many of the customers have churned?
plt.figure(figsize=(4,4))
sns.histplot(data=data, x='Churn')
plt.title('Customer churn')
plt.show()
```



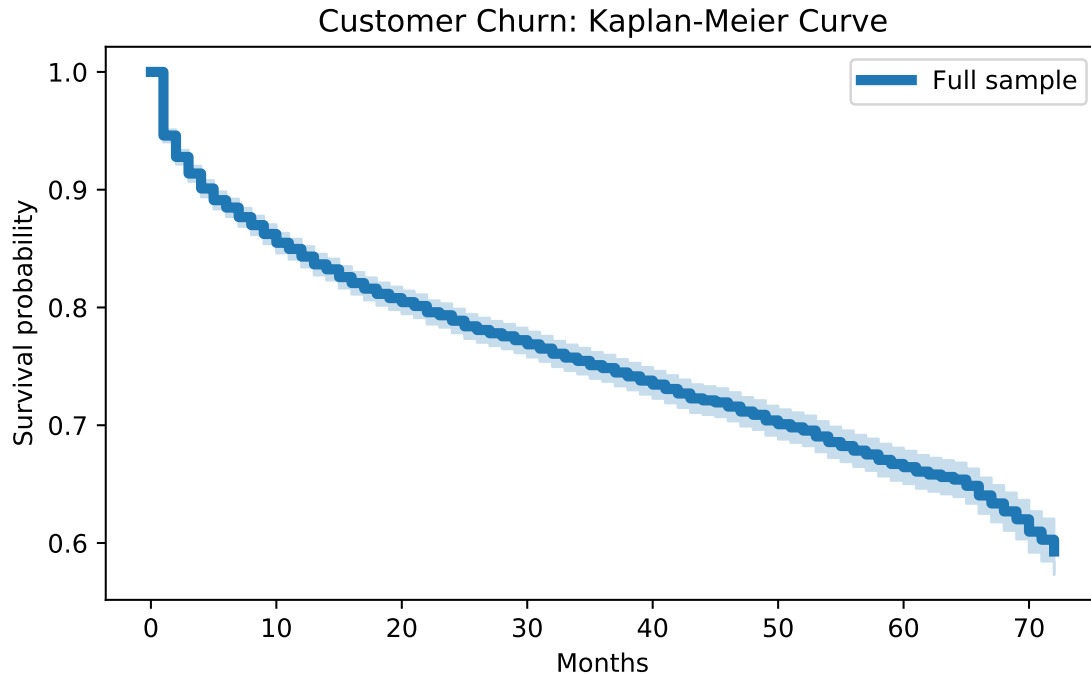
```
[8]: # How long on average were they/have they been a customer?
plt.figure(figsize=(4,4))
sns.barplot(data=data, x='Churn', y='tenure')
plt.title('Customer tenure')
plt.show()
```



The full sample – Kaplan-Meier curve

```
[9]: # i'll need to convert churn to a binary variable:  
data.replace({"Churn":{"No": 0, "Yes": 1}}, inplace=True)
```

```
[10]: kmf = KaplanMeierFitter()  
  
kmf.fit(data.tenure, data.Churn, label = 'Full sample')  
  
kmf.plot(linewidth=4, figsize=(7, 4))  
plt.title('Customer Churn: Kaplan-Meier Curve')  
plt.xlabel('Months')  
plt.ylabel('Survival probability')  
plt.show()
```



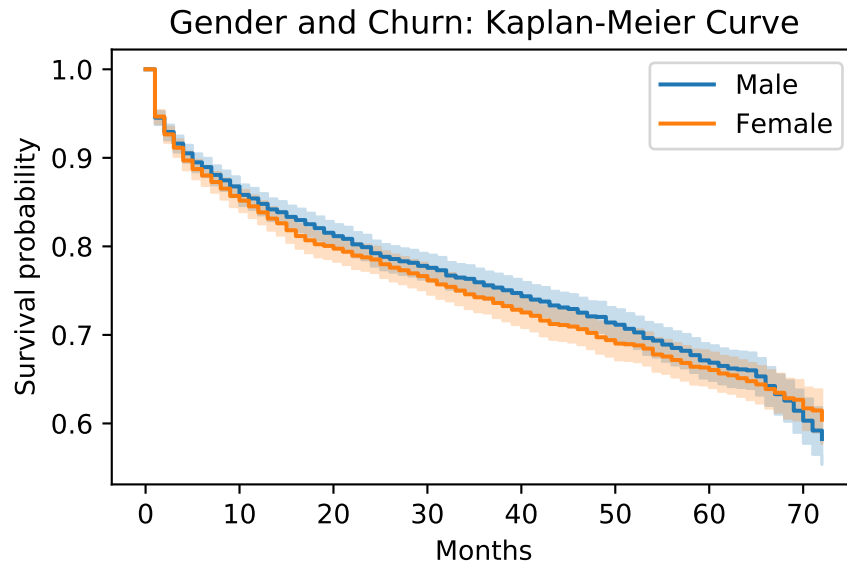
Next, I'm going to visualize the possible effect of a few variables on churn:

Males vs females

From the Kaplan-Meier curves below, gender doesn't appear to have much of an influence on customer churn.

```
[11]: df1 = data[data.gender == 'Male']
      df2 = data[data.gender == 'Female']

[12]: kmf.fit(df1.tenure, df1.Churn)
      kmf.plot(label = 'Male', figsize=(5,3))
      kmf.fit(df2.tenure, df2.Churn)
      kmf.plot(label = 'Female', figsize=(5,3))
      plt.title('Gender and Churn: Kaplan-Meier Curve')
      plt.xlabel('Months')
      plt.ylabel('Survival probability')
      plt.show()
```

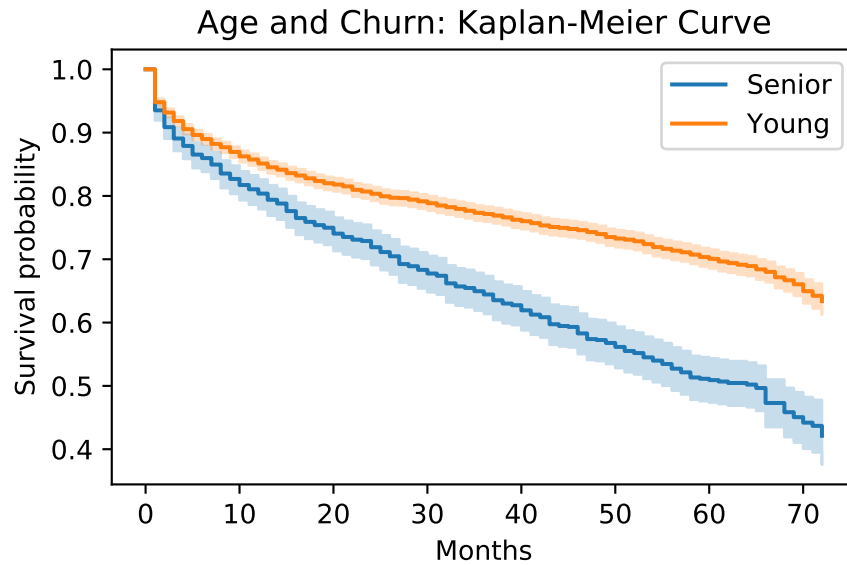


Senior citizen

Being a senior citizen has a large influence on customer churn, likely for multiple reasons!

```
[13]: df1 = data[data.SeniorCitizen == 1]
      df2 = data[data.SeniorCitizen == 0]

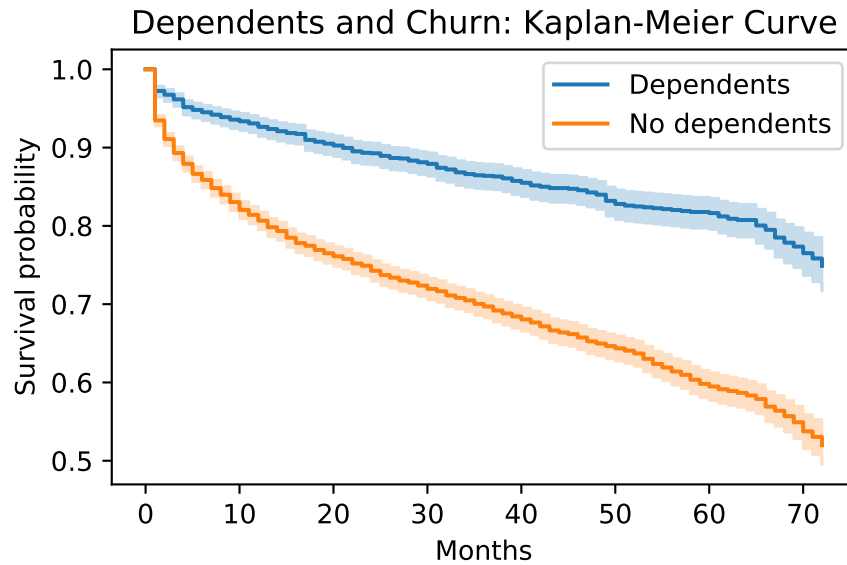
[14]: kmf.fit(df1.tenure, df1.Churn)
      kmf.plot(label = 'Senior', figsize=(5,3))
      kmf.fit(df2.tenure, df2.Churn)
      kmf.plot(label = 'Young', figsize=(5,3))
      plt.title('Age and Churn: Kaplan-Meier Curve')
      plt.xlabel('Months')
      plt.ylabel('Survival probability')
      plt.show()
```



Having dependents

Having dependents makes you much more likely to stay on as a customer.

```
[15]: df1 = data[data.Dependents == 'Yes']  
      df2 = data[data.Dependents == 'No']  
  
[16]: kmf.fit(df1.tenure, df1.Churn)  
      kmf.plot(label = 'Dependents', figsize=(5,3))  
      kmf.fit(df2.tenure, df2.Churn)  
      kmf.plot(label = 'No dependents', figsize=(5,3))  
      plt.title('Dependents and Churn: Kaplan-Meier Curve')  
      plt.xlabel('Months')  
      plt.ylabel('Survival probability')  
      plt.show()
```

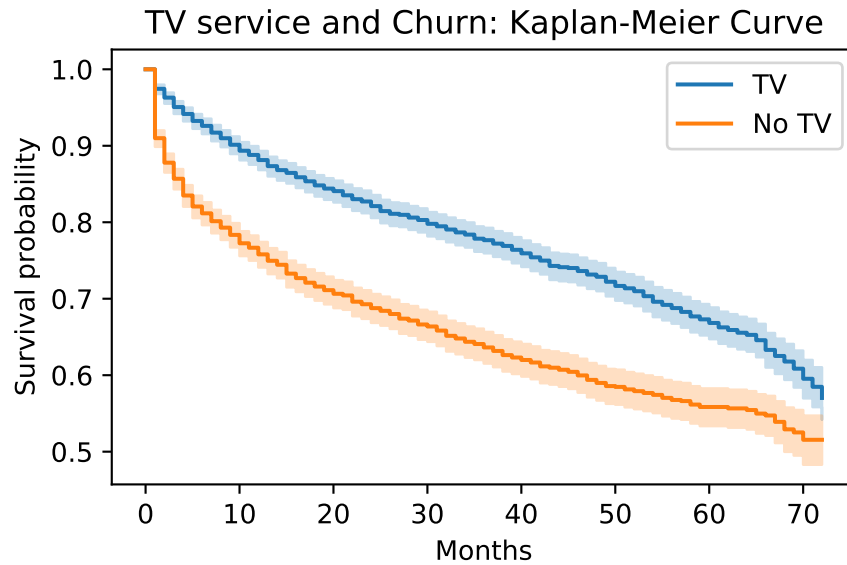



TV streaming

Finally, what about whether you pay for a for TV streaming? People who pay for TV streaming are less likely to churn, but that benefit does narrow slightly over time.

```
[17]: df1 = data[data.StreamingTV == 'Yes']
      df2 = data[data.StreamingTV == 'No']

[18]: kmf.fit(df1.tenure, df1.Churn)
      kmf.plot(label = 'TV', figsize=(5,3))
      kmf.fit(df2.tenure, df2.Churn)
      kmf.plot(label = 'No TV', figsize=(5,3))
      plt.title('TV service and Churn: Kaplan-Meier Curve')
      plt.xlabel('Months')
      plt.ylabel('Survival probability')
      plt.show()
```



1.4 CPH Models

To test the hazard rate of our features, I'm now going to fit and compare cox proportional hazard models.

Before building any models, let's take a look at the correlation between my features:

```
[19]: # dummy code all variables:
data1 = data[['Churn', 'gender', 'SeniorCitizen', 'Dependents', 'DeviceProtection', 'StreamingTV', 'StreamingMovies']]
data1.head()
```

```
[19]:   Churn  gender  SeniorCitizen  Dependents  DeviceProtection  StreamingTV  \
0      0  Female              0           No                No          No
1      0   Male              0           No                Yes          No
2      1   Male              0           No                No          No
3      0   Male              0           No                Yes          No
4      1  Female              0           No                No          No

   StreamingMovies
0              No
1              No
2              No
3              No
4              No
```

```
[20]:
```


```
data1 = pd.get_dummies(data1, drop_first=True).
↳drop(columns=['DeviceProtection_No internet service','StreamingTV_No_
↳internet service','StreamingMovies_No internet service'])
data1['months'] = data.tenure
data1['charges'] = zscore(data.MonthlyCharges)
data1.head()
```

```
[20]:
```

	Churn	SeniorCitizen	gender_Male	Dependents_Yes	DeviceProtection_Yes	\
0	0	0	0	0	0	
1	0	0	1	0	1	
2	1	0	1	0	0	
3	0	0	1	0	1	
4	1	0	0	0	0	

	StreamingTV_Yes	StreamingMovies_Yes	months	charges
0	0	0	1	-1.160323
1	0	0	34	-0.259629
2	0	0	2	-0.362660
3	0	0	45	-0.746535
4	0	0	2	0.197365

```
[21]: cors = data1.corr()
plt.figure(figsize=(5,4))
sns.heatmap(cors)
plt.show()
```



1.4.1 Model 1: Demographics

In this model, I'm just going to include age (senior or not), gender, and any dependents.

```
[22]: cph = CoxPHFitter()
cph.fit(data1,
        drop(columns=['charges', 'DeviceProtection_Yes', 'StreamingTV_Yes', 'StreamingMovies_Yes']),
        duration_col='months', event_col='Churn')
cph.print_summary(style='ascii')
```

```
<lifelines.CoxPHFitter: fitted with 7043 total observations, 5174 right-censored
observations>
```

```
    duration col = 'months'
    event col = 'Churn'
    baseline estimation = breslow
    number of observations = 7043
    number of events observed = 1869
    partial log-likelihood = -15502.35
    time fit was run = 2021-03-25 02:08:16 UTC
```


```

---
              coef exp(coef)  se(coef)  coef lower 95%  coef upper 95%
exp(coef) lower 95% exp(coef) upper 95%
covariate
SeniorCitizen    0.36      1.43    0.05          0.25          0.46
1.29              1.59
gender_Male     -0.03      0.97    0.05         -0.12          0.06
0.89              1.07
Dependents_Yes  -0.83      0.44    0.06         -0.95         -0.70
0.39              0.49

              z      p  -log2(p)
covariate
SeniorCitizen   6.60 <0.005    34.54
gender_Male    -0.59   0.55     0.85
Dependents_Yes -13.28 <0.005   131.20
---
Concordance = 0.60
Partial AIC = 31010.71
log-likelihood ratio test = 301.37 on 3 df
-log2(p) of ll-ratio test = 213.60

```

```
[23]: plt.figure(figsize=(4,2))
      cph.plot()
      plt.show()
```



As we could see from the above kaplan-meier plots, the first model confirms that age and dependents influence hazard risk over time.

1.4.2 Model 2: Demographics + Device protection + Charges

Same as prior model, now also accounting for whether customers pay for device protection and how much their monthly charges are

```
[24]: cph = CoxPHFitter()
cph.fit(data1.drop(columns=['StreamingTV_Yes', 'StreamingMovies_Yes']),
        duration_col='months', event_col='Churn')
cph.print_summary(style='ascii')
```

```
<lifelines.CoxPHFitter: fitted with 7043 total observations, 5174 right-censored
observations>
```

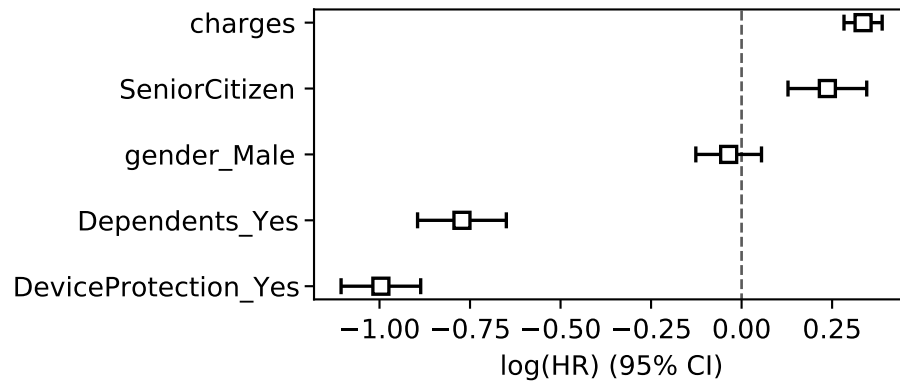
```
      duration col = 'months'
      event col = 'Churn'
      baseline estimation = breslow
      number of observations = 7043
      number of events observed = 1869
      partial log-likelihood = -15323.46
      time fit was run = 2021-03-25 02:08:17 UTC

---
      coef  exp(coef)  se(coef)  coef lower 95%  coef upper
95%  exp(coef) lower 95%  exp(coef) upper 95%
covariate
SeniorCitizen      0.24      1.27      0.06      0.13
0.35      1.14      1.41
gender_Male      -0.04      0.96      0.05      -0.13
0.05      0.88      1.06
Dependents_Yes     -0.77      0.46      0.06      -0.89
-0.65      0.41      0.52
DeviceProtection_Yes -1.00      0.37      0.06      -1.11
-0.89      0.33      0.41
charges           0.34      1.40      0.03      0.28
0.39      1.33      1.47
```

```
      z      p  -log2(p)
covariate
SeniorCitizen      4.27 <0.005      15.65
gender_Male      -0.77  0.44      1.18
Dependents_Yes     -12.35 <0.005     113.94
DeviceProtection_Yes -17.76 <0.005     232.01
charges           12.48 <0.005     116.29
```

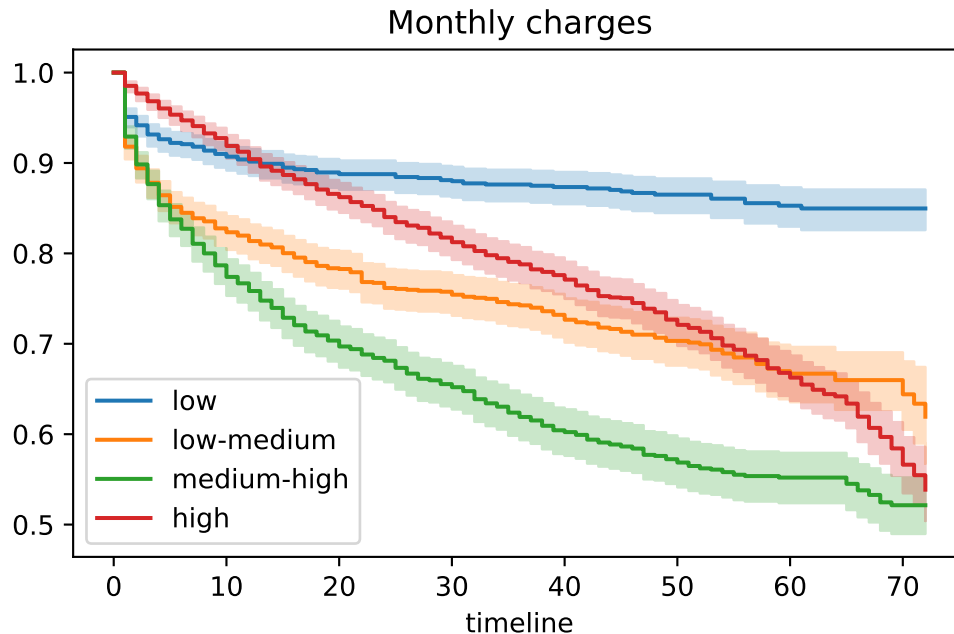
```
---
Concordance = 0.67
Partial AIC = 30656.92
log-likelihood ratio test = 659.16 on 5 df
-log2(p) of ll-ratio test = 463.34
```

```
[25]: plt.figure(figsize=(4,2))
      cph.plot()
      plt.show()
```



Plot the kaplan-meier curve at different levels of monthly charge:

```
[26]: qs = [0, 0.25, 0.5, 0.75, 1]
      labels = ['low', 'low-medium', 'medium-high', 'high']
      for q in range(len(qs)-1):
          low_q = data1.charges.quantile(qs[q])
          high_q = data1.charges.quantile(qs[q+1])
          this_data = data1[(data1.charges >= low_q) & (data1.charges < high_q)]
          kmf.fit(this_data.months, event_observed=this_data.Churn)
          kmf.plot(label=labels[q], figsize=(6,3.5))
      plt.title('Monthly charges')
      plt.show()
```



This model shows that customers who pay for device protection and are less likely to churn, and those who pay a higher monthly charge are more likely to churn. However, this relationship is pretty interesting – customers who have a high monthly charge are actually the least likely to churn within the first year, but then it becomes increasingly likely that they will churn. Customers who have a medium monthly charge are the most likely to churn, while customers who have a very low fee are the least likely to churn.

Overall, this model is a better fit to the data than the prior model as indicated by a lower AIC.

1.4.3 Model 3: All

Same as prior model, now also accounting for whether customers pay for TV and movie streaming

```
[27]: cph = CoxPHFitter()
      cph.fit(data1, duration_col='months', event_col='Churn')
      cph.print_summary(style='ascii')
```

```
<lifelines.CoxPHFitter: fitted with 7043 total observations, 5174 right-censored
observations>
```

```
      duration col = 'months'
      event col = 'Churn'
      baseline estimation = breslow
      number of observations = 7043
      number of events observed = 1869
      partial log-likelihood = -15282.64
      time fit was run = 2021-03-25 02:08:17 UTC
```

```
---
```


	coef	exp(coef)	se(coef)	coef lower 95%	coef upper 95%
covariate					
SeniorCitizen	0.22	1.25	0.06		0.11
	0.33	1.12	1.39		
gender_Male	-0.03	0.97	0.05		-0.13
	0.06	0.88	1.06		
Dependents_Yes	-0.75	0.47	0.06		-0.87
	-0.63	0.42	0.53		
DeviceProtection_Yes	-0.94	0.39	0.06		-1.05
	-0.83	0.35	0.44		
StreamingTV_Yes	-0.32	0.73	0.06		-0.44
	-0.20	0.65	0.82		
StreamingMovies_Yes	-0.38	0.69	0.06		-0.50
	-0.26	0.61	0.77		
charges	0.55	1.74	0.04		0.48
	0.62	1.62	1.86		

	z	p	-log2(p)
covariate			
SeniorCitizen	3.98	<0.005	13.81
gender_Male	-0.75	0.46	1.13
Dependents_Yes	-12.00	<0.005	107.85
DeviceProtection_Yes	-16.56	<0.005	202.23
StreamingTV_Yes	-5.26	<0.005	22.75
StreamingMovies_Yes	-6.25	<0.005	31.21
charges	15.15	<0.005	169.86


Concordance = 0.68

Partial AIC = 30579.28

log-likelihood ratio test = 740.80 on 7 df

-log2(p) of ll-ratio test = 514.76

```
[28]: plt.figure(figsize=(4,2))
      cph.plot()
      plt.show()
```



This final model shows that streaming TV or movies reduces the change of customer churn. This model is only marginally better than Model 2, as indicated by a slightly lower AIC.

1.5 Key Findings

The models show that having a lower monthly fee, having dependents, and having device protection are the main factor that mitigate customer churn. Accounting for more factors such as streaming services and age (senior citizen) also accounts for some variance in when customers churn, but gender is not an influential factor. The final model is likely to be the most appropriate to use.

1.6 Future directions and limitations

There are a number of ways in which these models could be improved. First, I did not account for any possible interactions between the features. For example, perhaps having dependents would reduce the negative impact of having a higher monthly charge rather than those features having an independent influence on customer churn. Second, data could be divided into training and test sets to test how well our features can predict customer churn in unseen data. Finally, more detailed analysis of customer demographics could be helpful, such as exact age, and also behavioral data such as the TV/movie viewing time, for example.

[]: